



# Airflow原理解析与常见坑答疑

— 基于v1.8.1

杭州优行科技-大数据部



汪涛

# CONTENTS

- 01** 基础介绍
- 02** 调度原理
- 03** 常用操作含义
- 04** 新增插件介绍
- 05** QA

01

## // 基础介绍



# 基础介绍 - 基础概念

**DAG:** Directed Acyclic Graph , is a collection of all the tasks you want to run, organized in a way that reflects their relationships and dependencies.

**DAG Run:** A DAG run is a physical instance of a DAG, containing task instances that run for a specific execution\_date

**execution\_date:** The execution\_date is the *logical* date and time which the DAG Run, and its task instances, are running for

**Operators:** While DAGs describe *how* to run a workflow, Operators determine what actually gets done.

**Sensor:** waits for a certain time, file, database row.

**tasks:** Once an operator is instantiated, it is referred to as a “task”. The instantiation defines specific values when calling the abstract operator, and the parameterized task becomes a node in a DAG

**Task Instances:** A task instance represents a specific run of a task

**Workflows:** By combining DAGs and Operators to create TaskInstances, you can build complex workflows

---

**Hooks:** Hooks are interfaces to external platforms and databases like Hive, S3, MySQL, Postgres, HDFS, and Pig.

**Pools:** Some systems can get overwhelmed when too many processes hit them at the same time, Airflow pools can be used to **limit the execution parallelism** on arbitrary sets of tasks

**Connections:** The connection information to external systems .

**Queues:** When using the CeleryExecutor, the Celery queues that tasks are sent to can be specified



配置文件 : airflow.cfg (\$AIRFLOW\_HOME目录下)

airflow webserver

airflow scheduler

airflow worker



# 基础介绍 - 基础组件

```
[admin@test31 ~]$ ps -ef | grep webserver | grep -v grep
admin  9404 19348 10 12:16 ?      00:00:16 [ready] gunicorn: worker [airflow-webserver]
admin 11026 19348 14 12:17 ?      00:00:16 [ready] gunicorn: worker [airflow-webserver]
admin 12867 19348 23 12:18 ?      00:00:16 [ready] gunicorn: worker [airflow-webserver]
admin 14641 19348 83 12:19 ?      00:00:16 [ready] gunicorn: worker [airflow-webserver]
admin 19100 1102 8 09:19 ?      00:14:35 /home/admin/program/airflow_venv/bin/python2 /home/admin/program/airflow_venv/bin/airflow webserver
admin 19348 19100 0 09:19 ?      00:00:06 gunicorn: master [airflow-webserver]

[admin@test31 ~]$ ps -ef | grep 'airflow sch' | grep -v grep
admin 17055 19073 0 12:20 ?      00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_ads/t_red_bag_car.py
admin 17056 19073 0 12:20 ?      00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/openwork/jobchain/dep_day_86.py
admin 17057 19073 0 12:20 ?      00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_ads/t_show_city_stat_1d_new.py
admin 17058 19073 0 12:20 ?      00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/bi/order/t_hermes_upms_weekly.py
admin 17059 19073 0 12:20 ?      00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_ads/t_caocao_recharge_pay_withdraw_1d.py
admin 17060 19073 0 12:20 ?      00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_ads/t_user_carbon_stat_insert.py
admin 17061 19073 0 12:20 ?      00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_model/demand-fcst/t_real_damend_10_30.py
admin 17062 19073 0 12:20 ?      00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_ods/cclog_recover_partitions.py
admin 19073 1102 6 12:05 ?      00:01:05 /home/admin/program/airflow_venv/bin/python2 /home/admin/program/airflow_venv/bin/airflow scheduler
```



```
[admin@test31 ~]$ ps -ef | grep celery | grep -v grep
admin  5356  1102  0 02:53 ?          00:00:27 [celeryd: celery@test31>MainProcess] -active- (worker)
admin  5435  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-1]
admin  5436  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-2]
admin  5437  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-3]
admin  5438  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-4]
admin  5439  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-5]
admin  5440  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-6]
admin  5441  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-7]
admin  5442  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-8]
admin  5443  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-9]
admin  5444  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-10]
admin  5445  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-11]
admin  5446  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-12]
admin  5447  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-13]
admin  5448  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-14]
admin  5449  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-15]
admin  5450  5356  0 02:53 ?          00:00:00 [celeryd: celery@test31:ForkPoolWorker-16]
```

02

## 调度原理



- 1、常见疑问
- 2、状态流转过程
- 3、调度核心类
- 4、scheduler运行过程
- 5、job运行过程
- 6、跨DAG依赖实现



- 1、重新运行一个任务，clear之后还需要点击run吗
- 2、一个dag中有a、b任务，重跑过去时间的任务时，为什么clear a之后b会运行
- 3、任务正在运行，点击clear或者mark success之后正在运行的任务会终止吗
- 4、新上线dag为什么在页面上看不到
- 5、新上线一个dag为什么一直跑历史任务
- 6、为什么有的任务达到超时时间没有停止运行，有时会同时跑两个任务



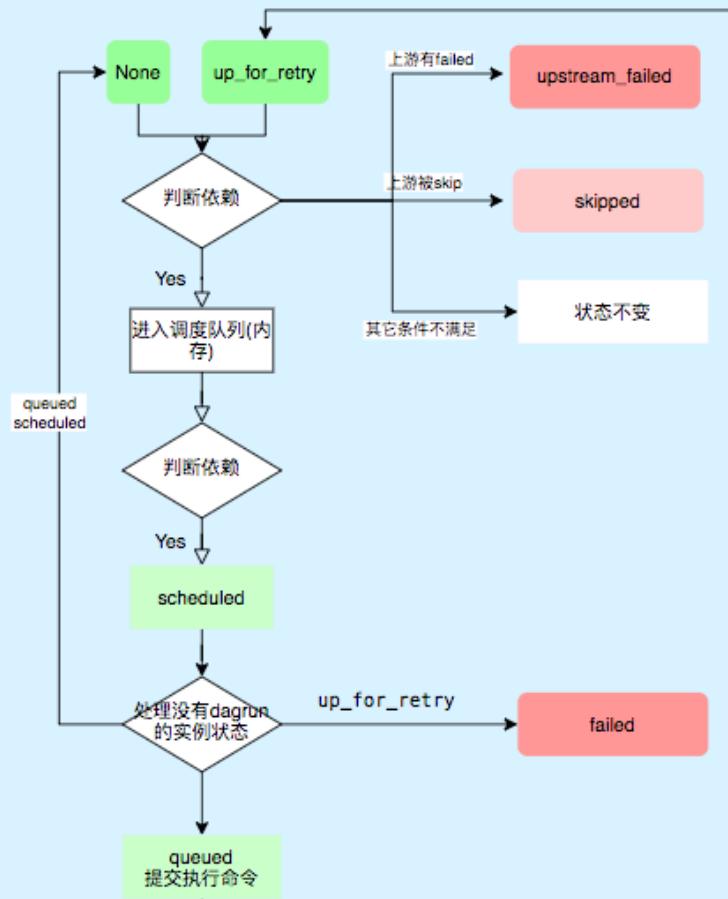
# 调度原理 - 状态流转

[■] success [■] running [■] failed [■] skipped [■] rescheduled [■] retry [■] queued [■] no status



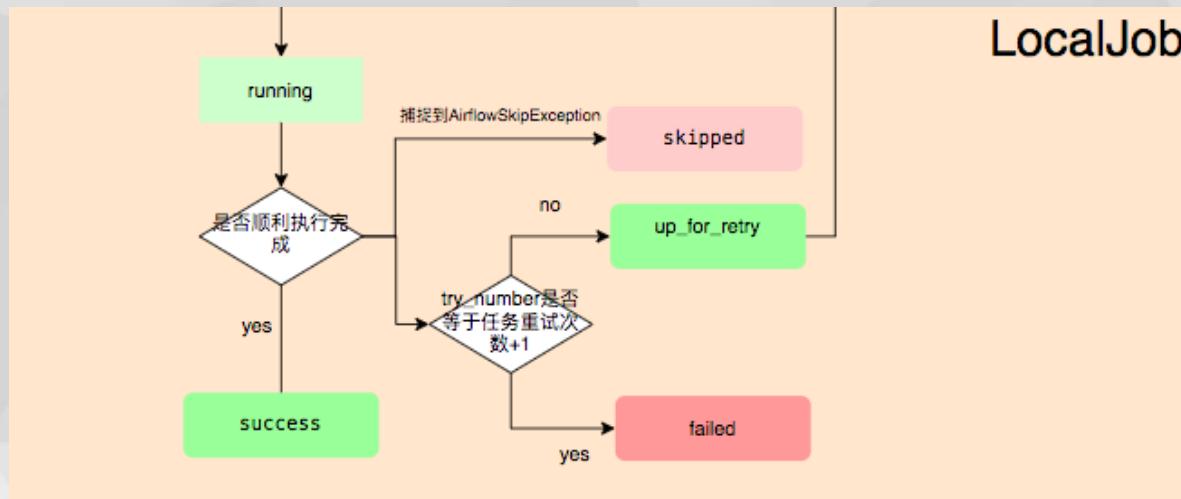
# 调度原理 - 状态流转

## SchedulerJob





# 调度原理 - 状态流转





SchedulerJob(BaseJob)

DagFileProcessManager

DagFileProcessor(AbstractDagFileProcessor)

CeleryExecutor(BaseExecutor)

DagRun

TaskInstance

LocalTaskJob(BaseJob)

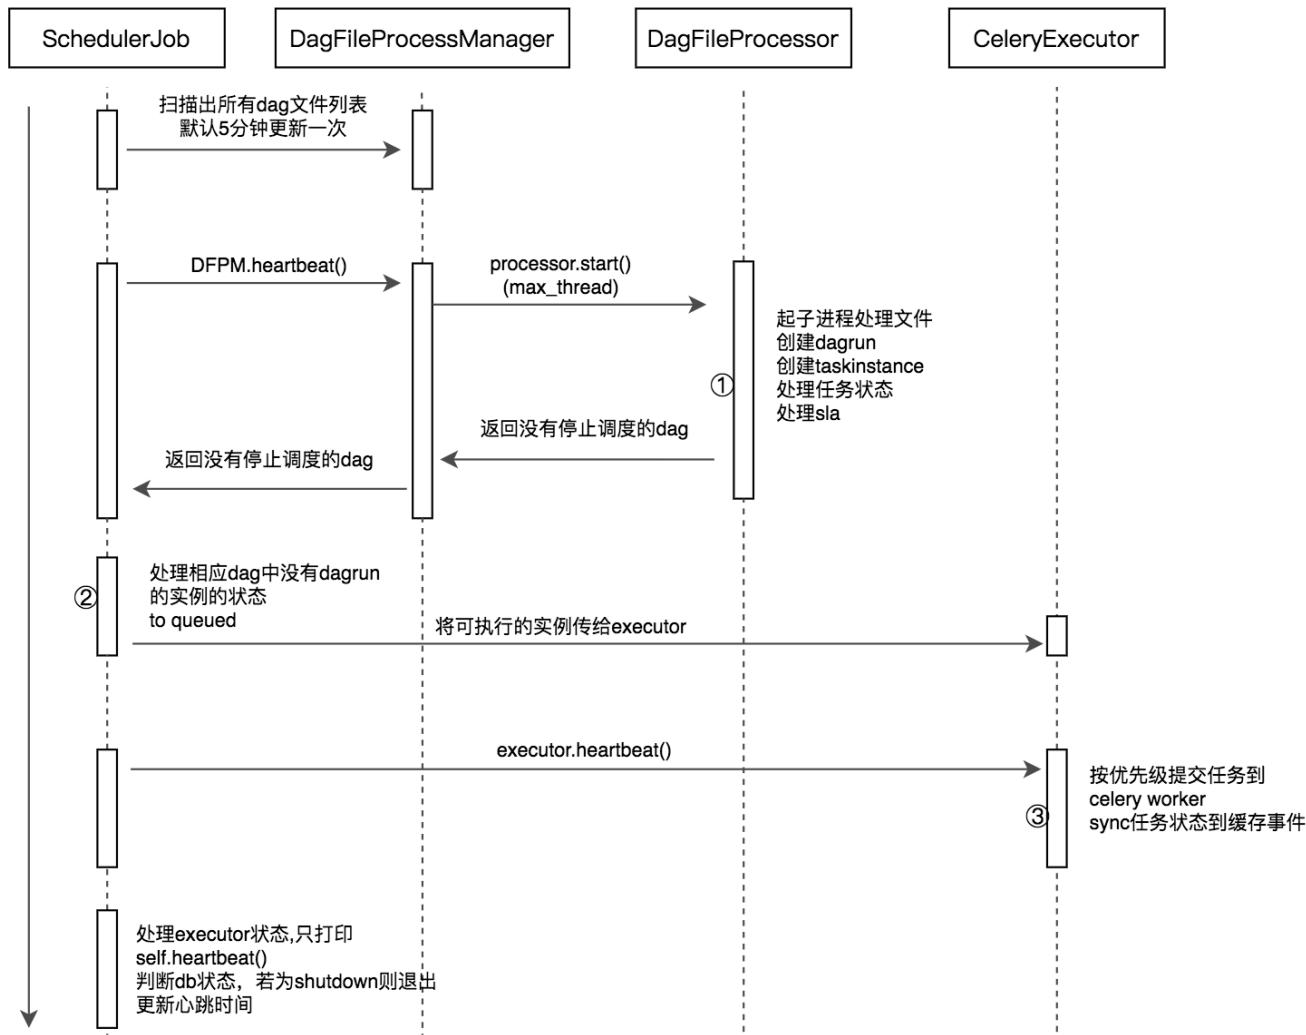
BashTaskRunner(BaseTaskRunner)

...Operator(BaseOperator)

.....



# 调度原理





# 调度原理

- scheduler过程

```
[operatoradmin@test31 ~]$ ps -ef | grep 'airflow sch'
admin    3817 1102 34 15:54 ?          00:00:03 /home/admin/program/airflow_venv/bin/python2 /home/admin/program/airflow_venv/bin/airflow scheduler
admin    4002 3817 99 15:54 ?          00:00:01 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_dw/dwdnew.py
admin    4018 3817 0 15:54 ?          00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/app/tags_finish_moiter.py
admin    4019 3817 0 15:54 ?          00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_ads/bi_new_gift_pack_stat_monthly.py
admin    4020 3817 0 15:54 ?          00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_ads/t_z_y_rush_order.py
admin    4021 3817 0 15:54 ?          00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/risk/operation_cclog_start_up.py
admin    4022 3817 0 15:54 ?          00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_ads/t_dws_cnt_interactive_count.py
admin    4023 3817 0 15:54 ?          00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/cc_ads/zhongyue_ads.py
admin    4024 3817 0 15:54 ?          00:00:00 airflow scheduler /home/admin/program/airflow/dags/python/dag/openwork/jobchain/dep_day_66.py
```

①

增加dag run

- 1、将超时的dag run状态置为failed
- 2、根据最新的dagrun的execution\_date以及是否catchup创建新的dag run

处理task instance

- 1、根据dag对所有在running中的dagrun增加和删除taskinstance，一次性补齐
- 2、根据dagrun所有taskinstance的状态更新dagrun的状态
- 3、将无状态和up\_for\_retry，判断依赖，如果上游有失败，状态置为upstream\_failed,上游有skipped状态置为skipped，依赖完成的任务放入调度队列

处理sla

将调度队列中的taskinstance再判断一次依赖，依赖完成后将taskinstance状态置为scheduled

处理僵尸任务：将job状态为running但ti不为running以及job心跳超时的ti标记为重试状态或失败状态，并调用相应的回调函数

返回脚本包含的没有停止调度的dag对象



# 调度原理

- scheduler过程

## create\_dag\_run

```
if not dag.catchup:  
    ...  
    now = datetime.now()  
    next_start = dag.following_schedule(now)  
    last_start = dag.previous_schedule(now)  
    if next_start <= now:  
        new_start = last_start  
    else:  
        new_start = dag.previous_schedule(last_start)  
  
    if dag.start_date:  
        if new_start >= dag.start_date:  
            dag.start_date = new_start  
        else:  
            dag.start_date = new_start  
  
if dag.schedule_interval == '@once':  
    period_end = next_run_date  
elif next_run_date:  
    period_end = dag.following_schedule(next_run_date)  
  
if next_run_date and period_end and period_end <= datetime.now():  
    next_run = dag.create_dagrun(  
        run_id='scheduled_' + next_run_date.isoformat(),  
        execution_date=next_run_date,  
        start_date=datetime.now(),  
        state=State.RUNNING,  
        external_trigger=False  
    )  
    return next_run
```



# 调度原理 - scheduler过程

## \_process\_task\_instances

```
active_dag_runs = []
for run in dag_runs:
    self.logger.info("Examining DAG run {}".format(run))
    # don't consider runs that are executed in the future
    if run.execution_date > datetime.now():
        self.logger.error("Execution date is in future: {}"
                           .format(run.execution_date))
        continue

    if len(active_dag_runs) >= dag.max_active_runs:
        self.logger.info("Active dag runs > max_active_run.")
        continue

    # skip backfill dagruns for now as long as they are not really sch
    if run.is_backfill:
        continue

    # todo: run.dag is transient but needs to be set
    run.dag = dag
    # todo: preferably the integrity check happens at dag collection t
    run.verify_integrity(session=session)
    run.update_state(session=session)
    if run.state == State.RUNNING:
        make_transient(run)
    active_dag_runs.append(run)
```

```
for run in active_dag_runs:
    self.logger.debug("Examining active DAG run {}".format(run))
    # this needs a fresh session sometimes tis get detached
    tis = run.get_task_instances(state=(State.NONE,
                                         State.UP_FOR_RETRY))

    ...
    for ti in tis:
        task = dag.get_task(ti.task_id)

        # fixme: ti.task is transient but needs to be set
        ti.task = task

        # future: remove adhoc
        if task.adhoc:
            continue

        if ti.are_dependencies_met(
            dep_context=DepContext(flag_upstream_failed=True),
            session=session):
            self.logger.debug('Queuing task: {}'.format(ti))
            queue.append(ti.key)

session.close()
```



# 调度原理 - scheduler过程

## 进入SCHEDULED状态

```
for ti_key in ti_keys_to_schedule:
    dag = dagbag.dags[ti_key[0]]
    task = dag.get_task(ti_key[1])
    ti = models.TaskInstance(task, ti_key[2])

    ti.refresh_from_db(session=session, lock_for_update=True)
    ...
    dep_context = DepContext(deps=QUEUE_DEPS, ignore_task_deps=True)
    ...
    if ti.are_dependencies_met(
        dep_context=dep_context,
        session=session,
        verbose=True):
        ...
        ti.state = State.SCHEDULED
        ...
# Also save this task instance to the DB.
self.logger.info("Creating / updating {} in ORM".format(ti))
session.merge(ti)
session.commit()
```

```
if flag_upstream_failed:
    if tr == TR.ALL_SUCCESS:
        if upstream_failed or failed:
            ti.set_state(State.UPSTREAM_FAILED, session)
    elif skipped:
        ti.set_state(State.SKIPPED, session)
    elif tr == TR.ALL_FAILED:
        if successes or skipped:
            ti.set_state(State.SKIPPED, session)
    elif tr == TR.ONE_SUCCESS:
        if upstream_done and not successes:
            ti.set_state(State.SKIPPED, session)
    elif tr == TR.ONE_FAILED:
        if upstream_done and not (failed or upstream_failed):
            ti.set_state(State.SKIPPED, session)
```



②

- 1、将dagrun不是running的taskinstance，retry状态置为failed，queued与scheduled置为None
- 2、将scheduled状态的taskinstance根据优先级和对应pool的slot数量(没有指定pool的默认128)将实例状态置为queued
- 3、将任务实例、优先级、队列(worker)以及执行命令发送到executor的queue(scheduler)中



# 调度原理 - scheduler过程

```
if len(simple_dags) > 0:  
    simple_dag_bag = SimpleDagBag(simple_dags)  
  
    ...  
    self._change_state_for_tis_without_dagrun(simple_dag_bag,  
                                              [State.UP_FOR_RETRY],  
                                              State.FAILED)  
    ...  
    self._change_state_for_tis_without_dagrun(simple_dag_bag,  
                                              [State.QUEUED,  
                                               State.SCHEDULED],  
                                              State.NONE)  
  
    self._execute_task_instances(simple_dag_bag,  
                               (State.SCHEDULED,))
```

```
current_task_concurrency = dag_id_to Possibly_running_task_count[dag_id]  
task_concurrency_limit = simple_dag_bag.get_dag(dag_id).concurrency  
self.logger.info("DAG {} has {} / {} running and queued tasks"  
                 .format(dag_id,  
                         current_task_concurrency,  
                         task_concurrency_limit))  
if current_task_concurrency >= task_concurrency_limit:  
    self.logger.info("Not executing {} since the number "  
                    "of tasks running or queued from DAG {}"  
                    " is >= to the "  
                    "DAG's task concurrency limit of {}"  
                    .format(task_instance,  
                            dag_id,  
                            task_concurrency_limit))  
    continue
```



# 调度原理

- scheduler 过程

```
command = " ".join(TI.generate_command(  
    task_instance.dag_id,  
    task_instance.task_id,  
    task_instance.execution_date,  
    local=True,  
    mark_success=False,  
    ignore_all_deps=False,  
    ignore_depends_on_past=False,  
    ignore_task_deps=False,  
    ignore_ti_state=False,  
    pool=task_instance.pool,  
    file_path=simple_dag_bag.get_dag(task_instance.dag_id).full_filepath,  
    pickle_id=simple_dag_bag.get_dag(task_instance.dag_id).pickle_id))  
  
priority = task_instance.priority_weight  
queue = task_instance.queue  
self.executor.queue_command(  
    task_instance,  
    command,  
    priority=priority,  
    queue=queue)
```



③

- 1、根据parallelism限制按照优先级将状态不为running的实例异步执行(提交任务到celery worker)
- 2、sync任务状态，将任务状态缓存到执行器的对象中(event\_buffer)



# 调度原理 - scheduler过程

```
def heartbeat(self):  
    # Triggering new jobs  
    if not self.parallelism:  
        open_slots = len(self.queued_tasks)  
    else:  
        open_slots = self.parallelism - len(self.running)  
  
    self.logger.debug("{} running task instances".format(len(self.running)))  
    self.logger.debug("{} in queue".format(len(self.queued_tasks)))  
    self.logger.debug("{} open slots".format(open_slots))  
  
    sorted_queue = sorted(  
        [(k, v) for k, v in self.queued_tasks.items()],  
        key=lambda x: x[1][1],  
        reverse=True)  
    for j in range(min((open_slots, len(self.queued_tasks))):  
        key, (command, _, queue, ti) = sorted_queue.pop(0)  
        ...  
        self.queued_tasks.pop(key)  
        ti.refresh_from_db()  
        if ti.state != State.RUNNING:  
            self.running[key] = command  
            self.execute_async(key, command=command, queue=queue)  
        else:  
            self.logger.debug(  
                'Task is already running, not sending to '  
                'executor: {}'.format(key))  
  
    # Calling child class sync method  
    self.logger.debug("Calling the {} sync method".format(self.__class__))  
    self.sync()
```



## 补充

### Start

---

生成SchedulerJob实例，置为running状态

扫描出dag文件列表

生成DagFileProcessorManager实例

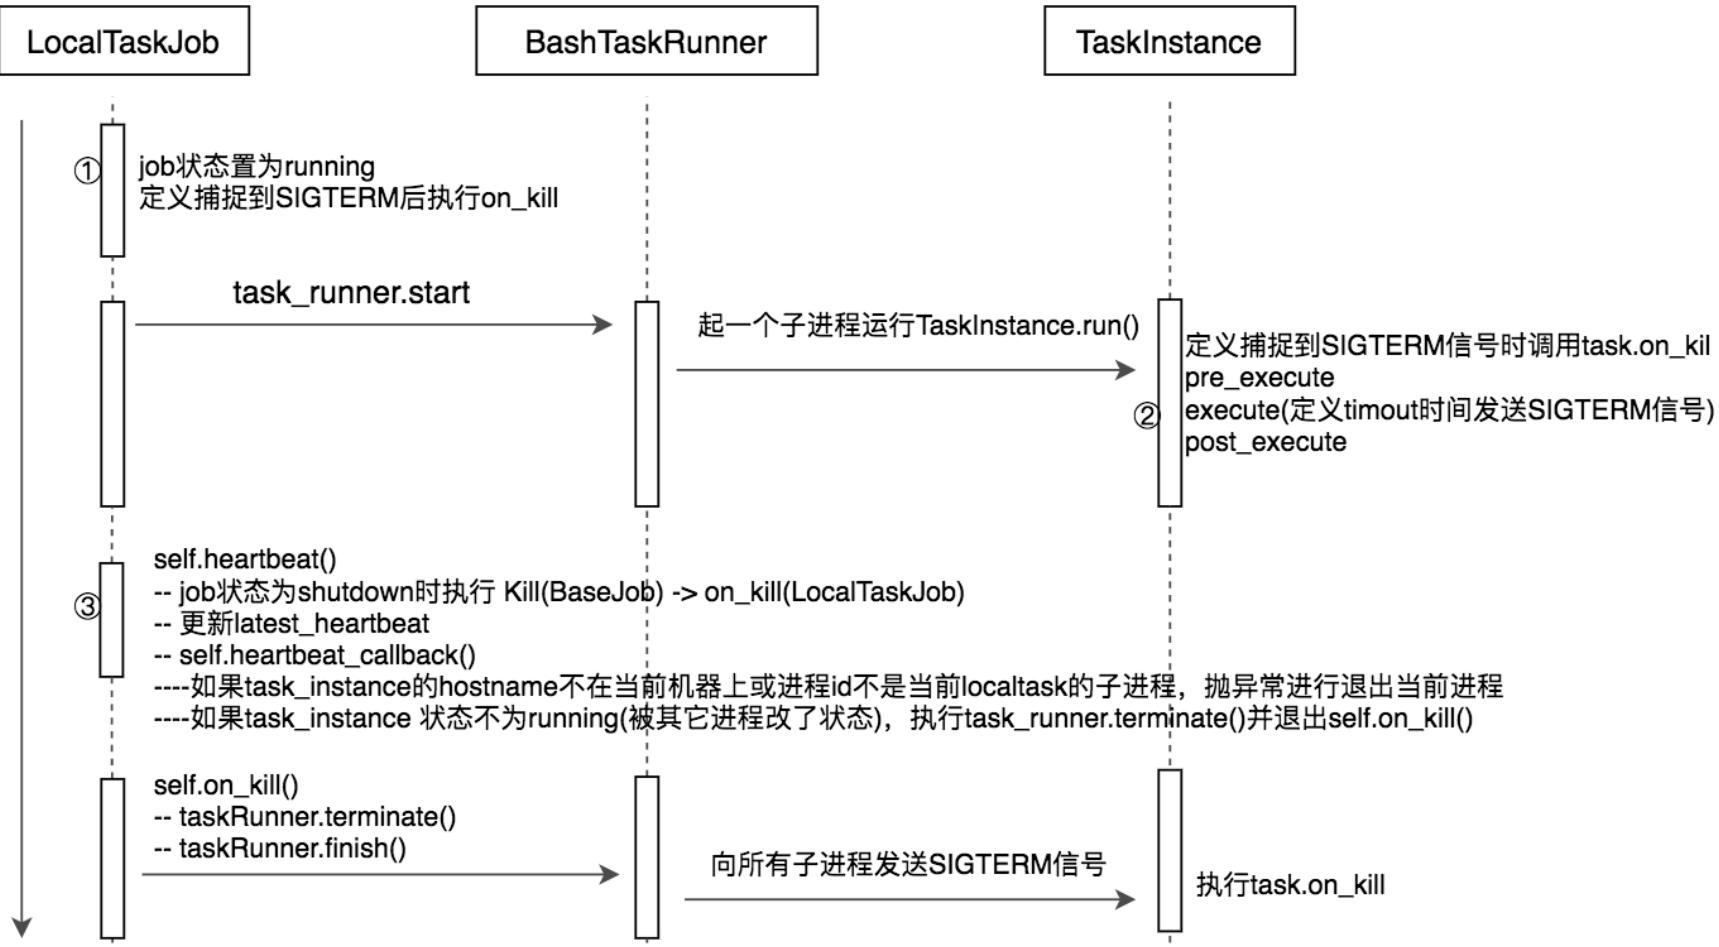
初始化Executor(CeleryExecutor)

### End

---

捕捉异常，kill掉所有子进程(先SIGTERM，再SIGKILL)

# 调度原理 - job执行过程





# 调度原理 - job执行过程

```
[admin@bd-spark03 ~]$ ps -ef | grep | grep 2913
admin 2913 2708 0 09:11 ? 00:00:02 /home/admin/program/airflow_venv/bin/python2 /home/admin/program/airflow_venv/bin/airflow run openwork.jobchain.dep_day_57 ningbo_yonghu_data 2020-04-20T09:10:00 --job_id 3978212 --raw -sd DAGS_FOLDER/python/dag/openwork/job
chain/dep_day_57.py
admin 3049 2913 0 09:11 ? 00:00:00 sh -c spark-sql --queue default --executor-memory 6G --num-executors 5 --executor-cores 3 --driver-memory 2G --hiveconf hive.cli.print.header=true -f /home/admin/program/airflow/dags/python/dag/openwork/job/sql/ningbo_yonghu_
data.sql --hiveconf current_date=2020-04-21 > /data1/data/bi/order/ningbo_yonghu_data/ningbo_yonghu_data2020-04-20.csv
[admin@bd-spark03 ~]$ ps -ef | grep | grep 2908
admin 2708 2913 0 09:11 ? 00:00:18 /home/admin/program/airflow_venv/bin/python2 /home/admin/program/airflow_venv/bin/airflow run openwork.jobchain.dep_day_57 ningbo_yonghu_data 2020-04-20T09:10:00 --local -sd /home/admin/program/airflow/dags/python/dag/openwork/job
rk/jobchain/dep_day_57.py
admin 2913 2708 0 09:11 ? 00:00:02 /home/admin/program/airflow_venv/bin/python2 /home/admin/program/airflow_venv/bin/airflow run openwork.jobchain.dep_day_57 ningbo_yonghu_data 2020-04-20T09:10:00 --job_id 3978212 --raw -sd DAGS_FOLDER/python/dag/openwork/job
chain/dep_day_57.py
[admin@bd-spark03 ~]$ ps -ef | grep | grep 9818
admin 2708 9818 1 09:11 ? 00:00:18 /home/admin/program/airflow_venv/bin/python2 /home/admin/program/airflow_venv/bin/airflow run openwork.jobchain.dep_day_57 ningbo_yonghu_data 2020-04-20T09:10:00 --local -sd /home/admin/program/airflow/dags/python/dag/openwo
rk/jobchain/dep_day_57.py
admin 9818 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-14]
[admin@bd-spark03 ~]$ ps -ef | grep 9776
root 29818 2 0 09:17 ? 00:00:00 [kworker/u:1]
admin 9818 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-1]
[admin@bd-spark03 ~]$ ps -ef | grep -v grep 9776
admin 8267 7972 0 05:36 ? 00:00:02 /home/admin/program/airflow_venv/bin/python2 /home/admin/program/airflow_venv/bin/airflow run bi.bi_cc_center_all.t_caoacao_city_operate_1d_part2 2020-04-20T04:27:00 --job_id 3977622 --raw -sd DAGS_FOLDER/python/dag/bi/bi_cc_
center_all.py
admin 9776 6434 0 Mar20 ? 04:05:24 [celeryd: celery@bd-spark03>MainProcess] -active- (worker)
admin 9791 9776 0 Mar20 ? 00:06:28 /home/admin/program/airflow_venv/bin/python2 /home/admin/program/airflow_venv/bin/airflow serve_logs
admin 9885 9776 0 Mar20 ? 00:00:19 [celeryd: celery@bd-spark03:ForkPoolWorker-1]
admin 9886 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-2]
admin 9887 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-3]
admin 9888 9776 0 Mar20 ? 00:00:19 [celeryd: celery@bd-spark03:ForkPoolWorker-4]
admin 9889 9776 0 Mar20 ? 00:00:19 [celeryd: celery@bd-spark03:ForkPoolWorker-5]
admin 9890 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-6]
admin 9811 9776 0 Mar20 ? 00:00:19 [celeryd: celery@bd-spark03:ForkPoolWorker-7]
admin 9812 9776 0 Mar20 ? 00:00:16 [celeryd: celery@bd-spark03:ForkPoolWorker-8]
admin 9813 9776 0 Mar20 ? 00:00:19 [celeryd: celery@bd-spark03:ForkPoolWorker-9]
admin 9814 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-10]
admin 9815 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-11]
admin 9816 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-12]
admin 9817 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-13]
admin 9818 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-14]
admin 9819 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-15]
admin 9820 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-16]
admin 9821 9776 0 Mar20 ? 00:00:19 [celeryd: celery@bd-spark03:ForkPoolWorker-17]
admin 9822 9776 0 Mar20 ? 00:00:18 [celeryd: celery@bd-spark03:ForkPoolWorker-18]
```



# 调度原理

-job执行过程

①

```
def on_kill(self):
    self.task_runner.terminate()
    self.task_runner.on_finish()

def terminate(self):
    if self.process and psutil.pid_exists(self.process.pid):
        kill_process_tree(self.logger, self.process.pid)

def on_finish(self):
    super(BashTaskRunner, self).on_finish()
```



②

## 1、先判断依赖，依赖没通过，退出执行

检查的依赖项有：实例要在以下状态中：scheduled, queued, failed, up\_for\_retry, upstream\_failed，以及一些可执行条件(执行日期超过当前时间、task超出有效期等)要求

2、taskinstance状态置为running,try\_number+1

3、初始化context

4、定义收到SIGTERM信号时调用task.on\_kill

5、清理taskinstance的XCom数据

6、渲染模板变量

7、执行task的pre\_execute

8、执行execute，若定义了timeout时间则发送一个定时器超时信SIGALRM，超时后执行task.on\_kill再抛异常退出当前进程  
(先定义了收到SIGALRM信号时的抛AirflowTaskTimeout异常，再捕捉AirflowTaskTimeout异常再执行task.on\_kill)

9、执行post\_execute

10、第3到第9步没有任务异常则将taskinstance状态置为success

11、第3到第9步有异常处理：

捕捉到AirflowSkipException时则将实例状态置为skipped

捕捉到其它异常则先在task\_fail中增加一条失败记录，如果实例的try\_number次数不等于配置的任务重试次数+1的倍数则将状态置为up\_for\_retry，否则置为failed(如果配置了发送邮件，则会将异常信息发送至任务指定的email中)，并调用重试或失败回调函数



# 调度原理 - job 执行过程

```
try:
    logging.info(msg.format(self=self))
    if not mark_success:
        context = self.get_template_context()

    task_copy = copy.copy(task)
    self.task = task_copy

    def signal_handler(signum, frame):
        """Setting kill signal handler"""
        logging.error("Killing subprocess")
        task_copy.on_kill()
        raise AirflowException("Task received SIGTERM signal")
    signal.signal(signal.SIGTERM, signal_handler)

    # Don't clear Xcom until the task is certain to execute
    self.clear_xcom_data()

    self.render_templates()
    task_copy.pre_execute(context=context)

    # If a timeout is specified for the task, make it fail
    # if it goes beyond
    result = None
    if task_copy.execution_timeout:
        try:
            with timeout(int(
                task_copy.execution_timeout.total_seconds())):
                result = task_copy.execute(context=context)
        except AirflowTaskTimeout:
            task_copy.on_kill()
            raise
    else:
        result = task_copy.execute(context=context)

    # If the task returns a result, push an XCom containing it
    if result is not None:
        self.xcom_push(key=XCOM_RETURN_KEY, value=result)

    task_copy.post_execute(context=context)
    Stats.incr('operator_successes_{}'.format(
        self.task.__class__.__name__), 1, 1)
    self.state = State.SUCCESS
except AirflowSkipException:
    self.state = State.SKIPPED
except (Exception, KeyboardInterrupt) as e:
    self.handle_failure(e, test_mode, context)
    raise
```

```
try:
    if task.retries and self.try_number % (task.retries + 1) != 0:
        self.state = State.UP_FOR_RETRY
        logging.info('Marking task as UP_FOR_RETRY')
        if task.email_on_retry and task.email:
            self.email_alert(error, is_retry=True)
    else:
        self.state = State.FAILED
        if task.retries:
            logging.info('All retries failed; marking task as FAILED')
        else:
            logging.info('Marking task as FAILED.')
            if task.email_on_failure and task.email:
                self.email_alert(error, is_retry=False)
    except Exception as e2:
        logging.error(
            'Failed to send email to: ' + str(task.email))
        logging.exception(e2)

    # Handling callbacks pessimistically
    try:
        if self.state == State.UP_FOR_RETRY and task.on_retry_callback:
            task.on_retry_callback(context)
        if self.state == State.FAILED and task.on_failure_callback:
            task.on_failure_callback(context)
    except Exception as e3:
        logging.error("Failed at executing callback")
        logging.exception(e3)

    if not test_mode:
        session.merge(self)
        session.commit()
        logging.error(str(error))
```



- 1、定义CrossDagDepPythonOperator(PythonOperator)
- 2、在pre\_execute实现：
  - 1) 记录该实例到表task\_instance\_ext中
  - 2) 判断依赖的其它dag任务是否成功，没有成功，则抛出AirflowSkipException
- 3、在post\_execute实现: 将task\_instance\_ext中实例状态置为success
- 4、起一个服务无限循环监控task\_instance\_ext表中状态，如果依赖的跨dag任务success，则clear该实例及其下游



# 调度原理 - 跨DAG依赖

```
is_task_ready = CrossDagDepPythonOperator.are_cross_dag_dep_conditions_met(
    super(CrossDagDepPythonOperator, self).dag_id,
    self.dag_dep_conditions, session, context['execution_date'])
if (not is_task_ready):
    ts_ext.state = STATE_WAITING
    session.merge(ts_ext)
    session.commit()
    raise AirflowSkipException("task dependencies are not ready yet!")

def post_execute(self, context):
    logging.info("post_execute in")
    session = settings.Session()
    execution_date = context['execution_date'].strftime('%Y-%m-%d %H:%M:%S')
    logging.info("execution_date:{}".format(execution_date))
    ts_ext = TaskInstanceExt(self.task_id, self.dag_id, execution_date)
    ts_ext.refresh_from_db(session)
    ts_ext.state = State.SUCCESS
    ts_ext.wait_time = (datetime.now() - ts_ext.start_time).total_seconds()
    session.merge(ts_ext)
    session.commit()
    session.close()

qry = session.query(TaskInstanceExt).filter(
    TaskInstanceExt.state == airflow_plugins.cross_dag_dep_python_operatior.STATE_WAITING,
    TaskInstanceExt.operator == CrossDagDepPythonOperator.__name__,
    TaskInstanceExt.execution_date > str_start_day,
    TaskInstanceExt.flag == 0,
    TaskInstanceExt.last_check_time < datetime.now() - timedelta(seconds=10),
)
```

03

## // 常用操作含义



# 常用操作含义

clear

- 1、如果不在RUNNING中直接删除
- 2、在RUNNING中将job置为shutdown
- 3、将dagrun置为running

```
def clear_task_instances(tis, session, activate_dag_runs=True):
    """
    Clears a set of task instances, but makes sure the running ones
    get killed.
    """

    job_ids = []
    for ti in tis:
        if ti.state == State.RUNNING:
            if ti.job_id:
                ti.state = State.SHUTDOWN
                job_ids.append(ti.job_id)
            # todo: this creates an issue with the webui tests
            # elif ti.state != State.REMOVED:
            #     ti.state = State.NONE
            #     session.merge(ti)
        else:
            session.delete(ti)
    if job_ids:
        from airflow.jobs import BaseJob as BJ
        for job in session.query(BJ).filter(BJ.id.in_(job_ids)).all():
            job.state = State.SHUTDOWN
    if activate_dag_runs:
        execution_dates = {ti.execution_date for ti in tis}
        dag_ids = {ti.dag_id for ti in tis}
        drs = session.query(DagRun).filter(
            DagRun.dag_id.in_(dag_ids),
            DagRun.execution_date.in_(execution_dates),
        ).all()
        for dr in drs:
            dr.state = State.RUNNING
            dr.start_date = datetime.now()
```



# 常用操作含义

run

依赖通过后传给executor执行

```
# Make sure the task instance can be queued
dep_context = DepContext(
    deps=QUEUE_DEPS,
    ignore_all_deps=ignore_all_deps,
    ignore_task_deps=ignore_task_deps,
    ignore_ti_state=ignore_ti_state)
failed_deps = list(ti.get_failed_dep_statuses(dep_context=dep_context))
if failed_deps:
    failed_deps_str = ", ".join(
        ["{}: {}".format(dep.dep_name, dep.reason) for dep in failed_deps])
    flash("Could not queue task instance for execution, dependencies not met: "
          "{}".format(failed_deps_str),
          "error")
    return redirect(origin)

executor.start()
executor.queue_task_instance(
    ti,
    ignore_all_deps=ignore_all_deps,
    ignore_task_deps=ignore_task_deps,
    ignore_ti_state=ignore_ti_state)
executor.heartbeat()
flash(
    "Sent {} to the message queue, "
    "it should start any moment now.".format(ti))
return redirect(origin)
```



# 常用操作含义

Mark success

- 1、task\_instance查漏补缺
- 2、将状态置为success

```
confirmed_dates = []
drs = DagRun.find(dag_id=dag.dag_id, execution_date=dates)
for dr in drs:
    dr.dag = dag
    dr.verify_integrity()
    confirmed_dates.append(dr.execution_date)
```

```
qry_dag = session.query(TI).filter(
    TI.dag_id==dag.dag_id,
    TI.execution_date.in_(confirmed_dates),
    TI.task_id.in_(task_ids)).filter(
        or_(TI.state.is_(None),
            TI.state != state)
    )
```

```
if commit:
    tis_altered = qry_dag.with_for_update().all()
    if len(sub_dag_ids) > 0:
        tis_altered += qry_sub_dag.with_for_update().all()
    for ti in tis_altered:
        ti.state = state
    session.commit()
```



- 1、重新运行一个任务，clear之后还需要点击run吗
- 2、一个dag中有a、b任务，重跑过去时间的任务时，为什么clear a之后b会运行
- 3、任务正在运行，点击clear或者mark success之后正在运行的任务会终止吗
- 4、新上线dag为什么在页面上看不到
- 5、新上线一个dag为什么一直跑历史任务
- 6、为什么有的任务达到超时时间还在跑



# 常用操作含义

- 常见疑问

```
class PythonOperator(BaseOperator):
    """Executes a Python callable..."""
    template_fields = ('templates_dict',)
    template_ext = tuple()
    ui_color = '#ffefeb'

    @apply_defaults
    def __init__(
            self,
            python_callable,
            op_args=None,
            op_kwargs=None,
            provide_context=False,
            templates_dict=None,
            templates_exts=None,
            *args, **kwargs):
        ...

    def execute(self, context):...
```

```
class CCPythonOperator(PythonOperator):

    def on_kill(self):
        import psutil
        p = psutil.Process(os.getpid())
        children = [x for x in
                   p.children(recursive=True)
                   if x.is_running()]
        for child in children:
            if child.is_running():
                child.terminate()
        for child in children:
            if child.is_running():
                child.kill()
```

```
from airflow.operators.python_operator import PythonOperator
to
from airflow_plugins.python_operator import CCPythonOperator as PythonOperator
```

04

## // 新增插件介绍



# 常用操作含义

- 新增插件介绍

```
class HiveToMysqlByDaxxOperator(BaseOperator):  
    def __init__(self):  
        pass  
  
    def execute(self, context):  
        pass  
  
    def on_kill(self):  
        u"""如果会产生子进程则应实现此方法清理子进程"""  
        pass
```



# 常用操作含义 - 新增插件介绍

```
class HiveToMysqlByDataxOperator(BaseOperator):  
    ui_color = '#eddede'  
    template_fields = [  
        'hive_table',  
        'partition_path',  
        'mysql_table',  
        'columns_map',  
        'clear_where_condition',  
        'write_mode'  
    ]  
  
    job_text = """..."""  
  
    def __init__(self,  
                 hive_table,  
                 partition_path,  
                 mysql_table,  
                 mysql_conn,  
                 columns_map,  
                 clear_where_condition=None,  
                 write_mode='insert',  
                 channel=1,  
                 datax_conn_id='datax_default',  
                 hive_metastore_conn_id='metastore_default',  
                 hdfs_file_encoding='utf-8',  
                 null_format=None,  
                 compress=None,  
                 jdbc_url_suffix="characterEncoding=utf-8",  
                 datax_jvm_para=None,  
                 datax_run_mode=None,  
                 datax_loglevel=None,  
                 datax_debug=None,  
                 error_limit_record=0,  
                 error_limit_percentage=0.0,  
                 *args,  
                 **kwargs):...  
  
    def execute(self, context):...  
  
    def on_kill(self):...
```

05

// QA

*Thanks!*